

PROJECT - OCD Patient Dataset: Demographics & Clinical Data

```
import pandas as pd

# Load the dataset
file_path = '/content/OCD Patient Dataset_ Demographics & Clinical Data.csv'
df = pd.read_csv(file_path)

# Dataset overview
print("Dataset Shape:", df.shape)
print("\nColumns:", df.columns.tolist())
print("\nFirst 5 rows:")
print(df.head())

# Check missing values
print("\nMissing values per column:")
print(df.isnull().sum())

Dataset Shape: (1500, 17)

Columns: ['Patient ID', 'Age', 'Gender', 'Ethnicity', 'Marital Status',
'Education Level', 'OCD Diagnosis Date', 'Duration of Symptoms (months)',
'Previous Diagnoses', 'Family History of OCD', 'Obsession Type', 'Compulsion
Type', 'Y-BOCS Score (Obsessions)', 'Y-BOCS Score (Compulsions)', 'Depression
Diagnosis', 'Anxiety Diagnosis', 'Medications']

First 5 rows:
   Patient ID  Age  Gender Ethnicity Marital Status Education Level \
0         1018  32  Female    African        Single    Some College
1         2406  69     Male    African       Divorced    Some College
2         1188  57     Male   Hispanic       Divorced  College Degree
3         6200  27  Female   Hispanic      Married  College Degree
4         5824  56  Female   Hispanic      Married     High School

   OCD Diagnosis Date Duration of Symptoms (months) Previous Diagnoses \
0           2016-07-15                  203                      MDD
1           2017-04-28                  180                      NaN
2           2018-02-02                  173                      MDD
3           2014-08-25                  126                     PTSD
4           2022-02-20                  168                     PTSD

   Family History of OCD Obsession Type Compulsion Type \
0                 No    Harm-related       Checking
1                Yes    Harm-related       Washing
2                 No  Contamination       Checking
3                Yes      Symmetry       Washing
4                 Yes     Hoarding       Ordering
```

	Y-BOCS Score (Obsessions)	Y-BOCS Score (Compulsions)	Depression Diagnosis
\			
0	17	10	Yes
1	21	25	Yes
2	3	4	No
3	14	28	Yes
4	39	18	No

	Anxiety Diagnosis	Medications
0	Yes	SNRI
1	Yes	SSRI
2	No	Benzodiazepine
3	Yes	SSRI
4	No	NaN

Missing values per column:

Patient ID	0
Age	0
Gender	0
Ethnicity	0
Marital Status	0
Education Level	0
OCD Diagnosis Date	0
Duration of Symptoms (months)	0
Previous Diagnoses	248
Family History of OCD	0
Obsession Type	0
Compulsion Type	0
Y-BOCS Score (Obsessions)	0
Y-BOCS Score (Compulsions)	0
Depression Diagnosis	0
Anxiety Diagnosis	0
Medications	386

dtype: int64

```
# Fill missing values
df['Previous Diagnoses'] = df['Previous Diagnoses'].fillna('Unknown')
df['Medications'] = df['Medications'].fillna('Unknown')

# Encode categorical columns
df['Gender'] = df['Gender'].replace({'Female':1, 'Male':2})
df['Ethnicity'] = df['Ethnicity'].replace({'African':1, 'Hispanic':2,
'Asian':3, 'Caucasian':4})
df['Marital Status'] = df['Marital Status'].replace({'Single':1,
'Divorced':2, 'Married':3})
df['Education Level'] = df['Education Level'].replace({'Some College':1,
'College Degree':2, 'High School':3, 'Graduate Degree':4})
df['Family History of OCD'] = df['Family History of OCD'].replace({'No':1,
'Yes':2})
df['Obsession Type'] = df['Obsession Type'].replace({'Harm-related':1,
```

```

'Contamination':2, 'Symmetry':3, 'Hoardings':4, 'Religious':5})
df['Compulsion Type'] = df['Compulsion Type'].replace({'Checking':1,
'Washing':2, 'Ordering':3, 'Praying':4, 'Counting':5})
df['Depression Diagnosis'] = df['Depression Diagnosis'].replace({'No':1,
'Yes':2})
df['Anxiety Diagnosis'] = df['Anxiety Diagnosis'].replace({'No':1, 'Yes':2})
df['Medications'] = df['Medications'].replace({'SNRI':0, 'SSRI':1,
'Benzodiazepine':2, 'Unknown':3})

# Drop OCD Diagnosis Date
df = df.drop(columns=['OCD Diagnosis Date'])

# Verify cleaning
print(df.info())
print(df.isnull().sum())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Patient ID      1500 non-null    int64  
 1   Age              1500 non-null    int64  
 2   Gender            1500 non-null    int64  
 3   Ethnicity         1500 non-null    int64  
 4   Marital Status   1500 non-null    int64  
 5   Education Level  1500 non-null    int64  
 6   Duration of Symptoms (months) 1500 non-null    int64  
 7   Previous Diagnoses 1500 non-null    object  
 8   Family History of OCD 1500 non-null    int64  
 9   Obsession Type   1500 non-null    int64  
 10  Compulsion Type  1500 non-null    int64  
 11  Y-BOCS Score (Obsessions) 1500 non-null    int64  
 12  Y-BOCS Score (Compulsions) 1500 non-null    int64  
 13  Depression Diagnosis 1500 non-null    int64  
 14  Anxiety Diagnosis   1500 non-null    int64  
 15  Medications        1500 non-null    int64  
dtypes: int64(15), object(1)
memory usage: 187.6+ KB
None
Patient ID          0
Age                 0
Gender               0
Ethnicity            0
Marital Status       0
Education Level      0
Duration of Symptoms (months) 0
Previous Diagnoses  0
Family History of OCD 0
Obsession Type       0

```

```
Compulsion Type          0
Y-BOCS Score (Obsessions) 0
Y-BOCS Score (Compulsions) 0
Depression Diagnosis      0
Anxiety Diagnosis         0
Medications                 0
dtype: int64

/tmp/ipython-input-319515668.py:6: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Gender'] = df['Gender'].replace({'Female':1, 'Male':2})
/tmp/ipython-input-319515668.py:7: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Ethnicity'] = df['Ethnicity'].replace({'African':1, 'Hispanic':2,
'Asian':3, 'Caucasian':4})
/tmp/ipython-input-319515668.py:8: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Marital Status'] = df['Marital Status'].replace({'Single':1,
'Divorced':2, 'Married':3})
/tmp/ipython-input-319515668.py:9: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Education Level'] = df['Education Level'].replace({'Some College':1,
'College Degree':2, 'High School':3, 'Graduate Degree':4})
/tmp/ipython-input-319515668.py:10: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Family History of OCD'] = df['Family History of OCD'].replace({'No':1,
'Yes':2})
/tmp/ipython-input-319515668.py:11: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Obsession Type'] = df['Obsession Type'].replace({'Harm-related':1,
'Contamination':2, 'Symmetry':3, 'Hoarding':4, 'Religious':5})
/tmp/ipython-input-319515668.py:12: FutureWarning: Downcasting behavior in
```

```
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Compulsion Type'] = df['Compulsion Type'].replace({'Checking':1,
'Washing':2, 'Ordering':3, 'Praying':4, 'Counting':5})
/tmp/ipython-input-319515668.py:13: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Depression Diagnosis'] = df['Depression Diagnosis'].replace({'No':1,
'Yes':2})
/tmp/ipython-input-319515668.py:14: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Anxiety Diagnosis'] = df['Anxiety Diagnosis'].replace({'No':1,
'Yes':2})
/tmp/ipython-input-319515668.py:15: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain
the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
    df['Medications'] = df['Medications'].replace({'SNRI':0, 'SSRI':1,
'Benzodiazepine':2, 'Unknown':3})

import matplotlib.pyplot as plt
import seaborn as sns

# 1. Demographics
plt.figure(figsize=(10,5))
sns.histplot(df['Age'], bins=20, kde=True)
plt.title('Age Distribution')
plt.show()

sns.countplot(x='Gender', data=df)
plt.title('Gender Distribution')
plt.show()

sns.countplot(x='Ethnicity', data=df)
plt.title('Ethnicity Distribution')
plt.show()

sns.countplot(x='Marital Status', data=df)
plt.title('Marital Status Distribution')
plt.show()

sns.countplot(x='Education Level', data=df)
```

```

plt.title('Education Level Distribution')
plt.show()

# 2. Clinical Data
sns.histplot(df['Duration of Symptoms (months)'], bins=20, kde=True)
plt.title('Duration of Symptoms')
plt.show()

sns.histplot(df['Y-BOCS Score (Obsessions)'], bins=20, kde=True)
plt.title('Y-BOCS Score (Obsessions)')
plt.show()

sns.histplot(df['Y-BOCS Score (Compulsions)'], bins=20, kde=True)
plt.title('Y-BOCS Score (Compulsions)')
plt.show()

sns.countplot(x='Depression Diagnosis', data=df)
plt.title('Depression Diagnosis Distribution')
plt.show()

sns.countplot(x='Anxiety Diagnosis', data=df)
plt.title('Anxiety Diagnosis Distribution')
plt.show()

# 3. OCD Symptom Types
sns.countplot(x='Obsession Type', data=df)
plt.title('Obsession Type Distribution')
plt.show()

sns.countplot(x='Compulsion Type', data=df)
plt.title('Compulsion Type Distribution')
plt.show()

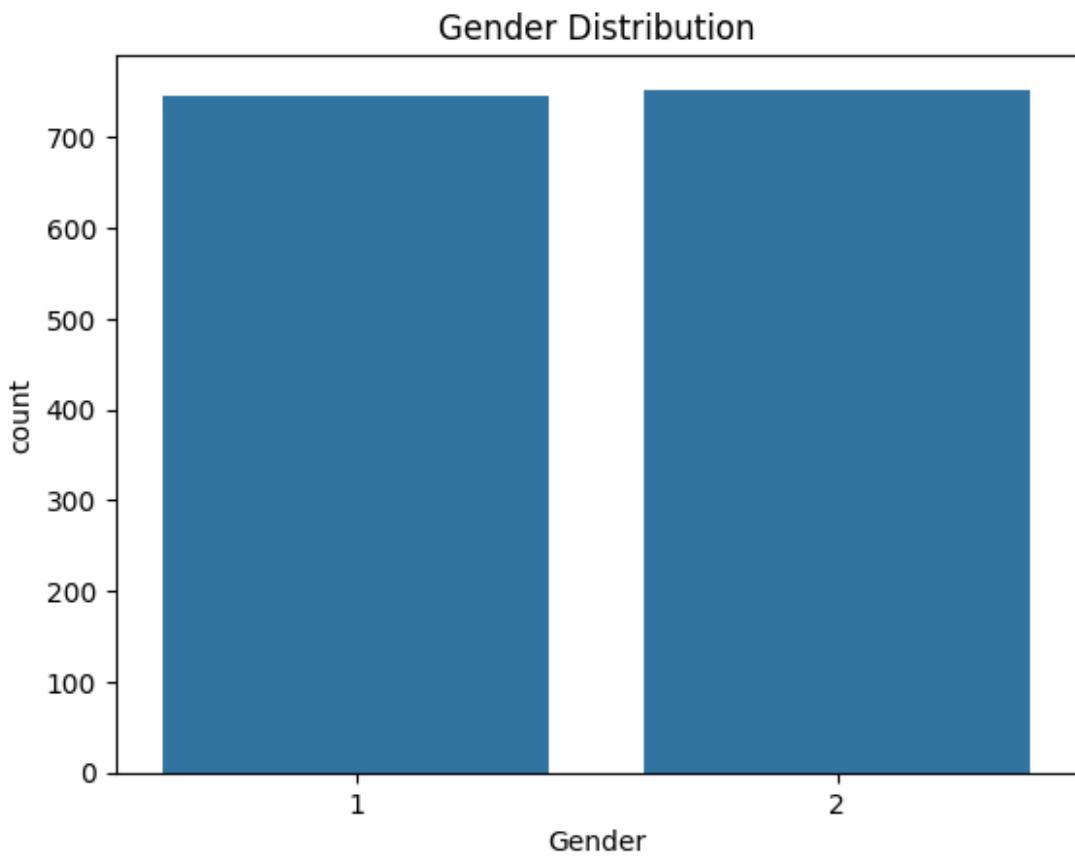
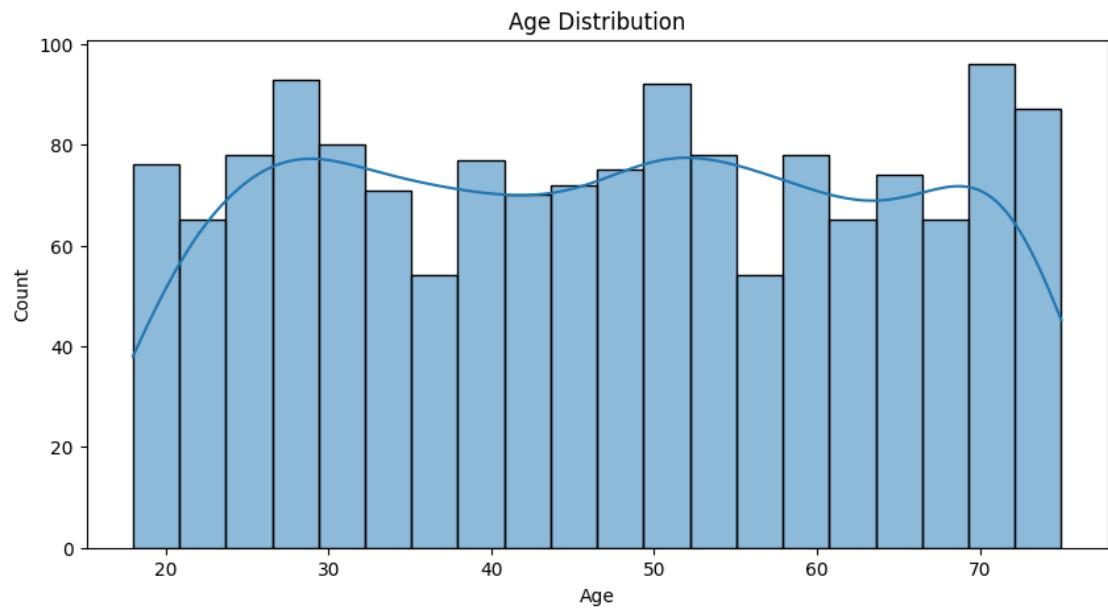
sns.countplot(x='Family History of OCD', data=df)
plt.title('Family History of OCD')
plt.show()

# 4. Medications
sns.countplot(x='Medications', data=df)
plt.title('Medication Type Distribution')
plt.show()

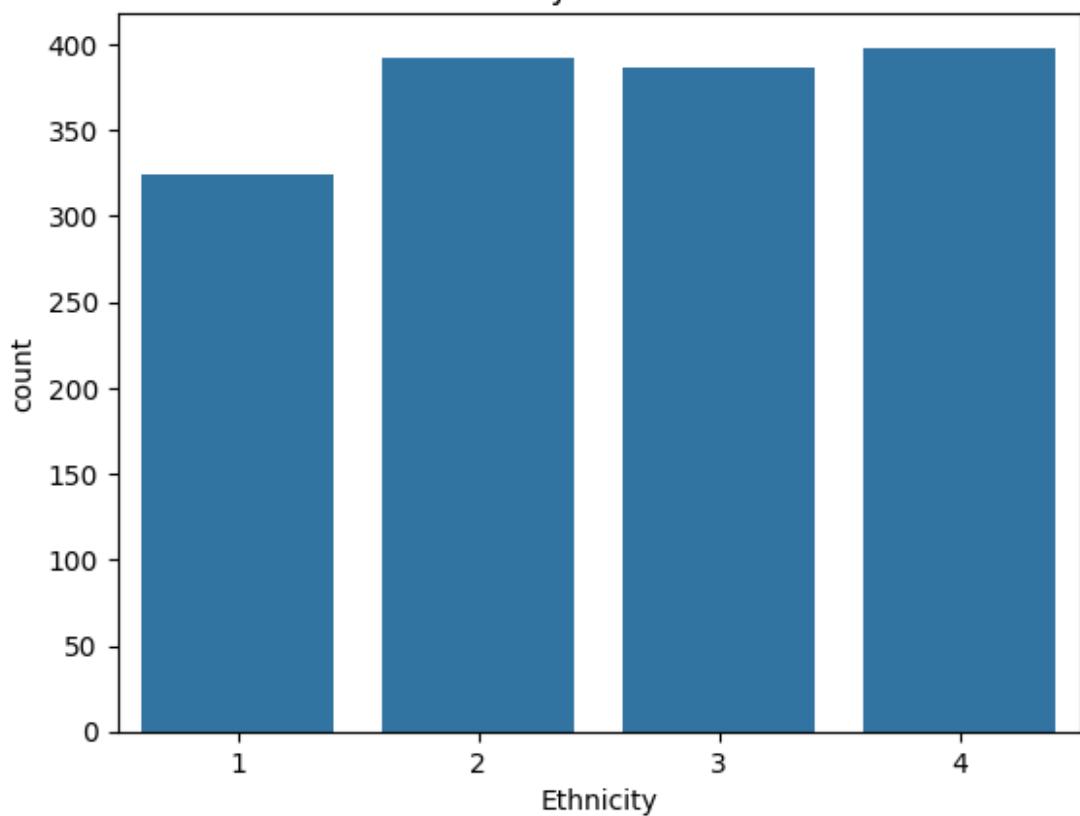
# 5. Correlation Matrix
numeric_cols = ['Age', 'Duration of Symptoms (months)', 'Obsession Type',
'Compulsion Type',
'Y-BOCS Score (Obsessions)', 'Y-BOCS Score (Compulsions)',
'Depression Diagnosis', 'Anxiety Diagnosis', 'Medications']
plt.figure(figsize=(12,10))
sns.heatmap(df[numeric_cols].corr(), annot=True, cmap='coolwarm', fmt=".2f")

```

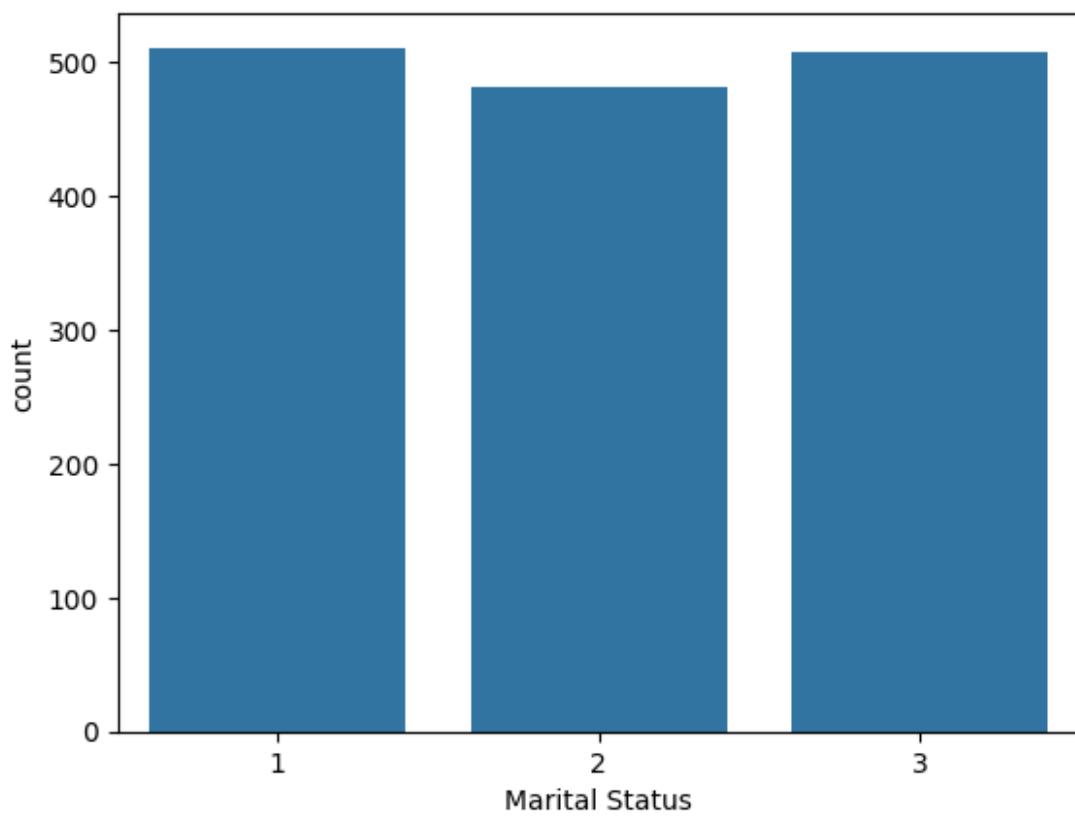
```
plt.title('Correlation Matrix')  
plt.show()
```



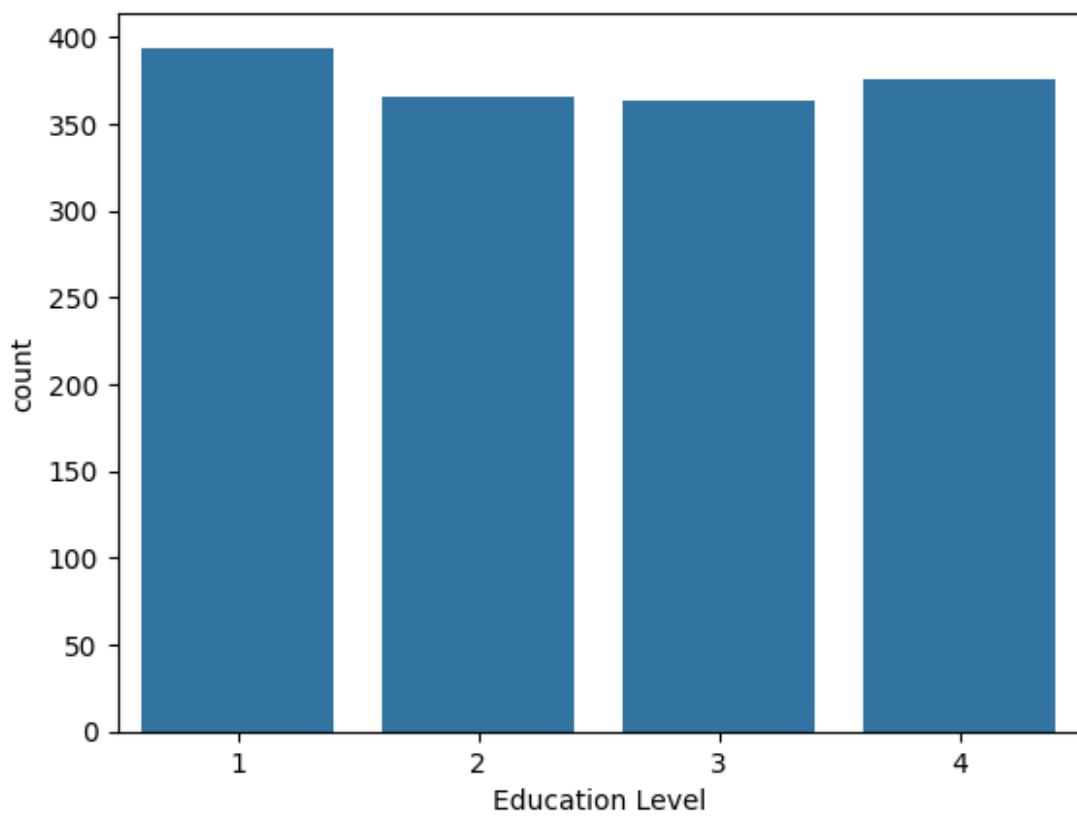
Ethnicity Distribution

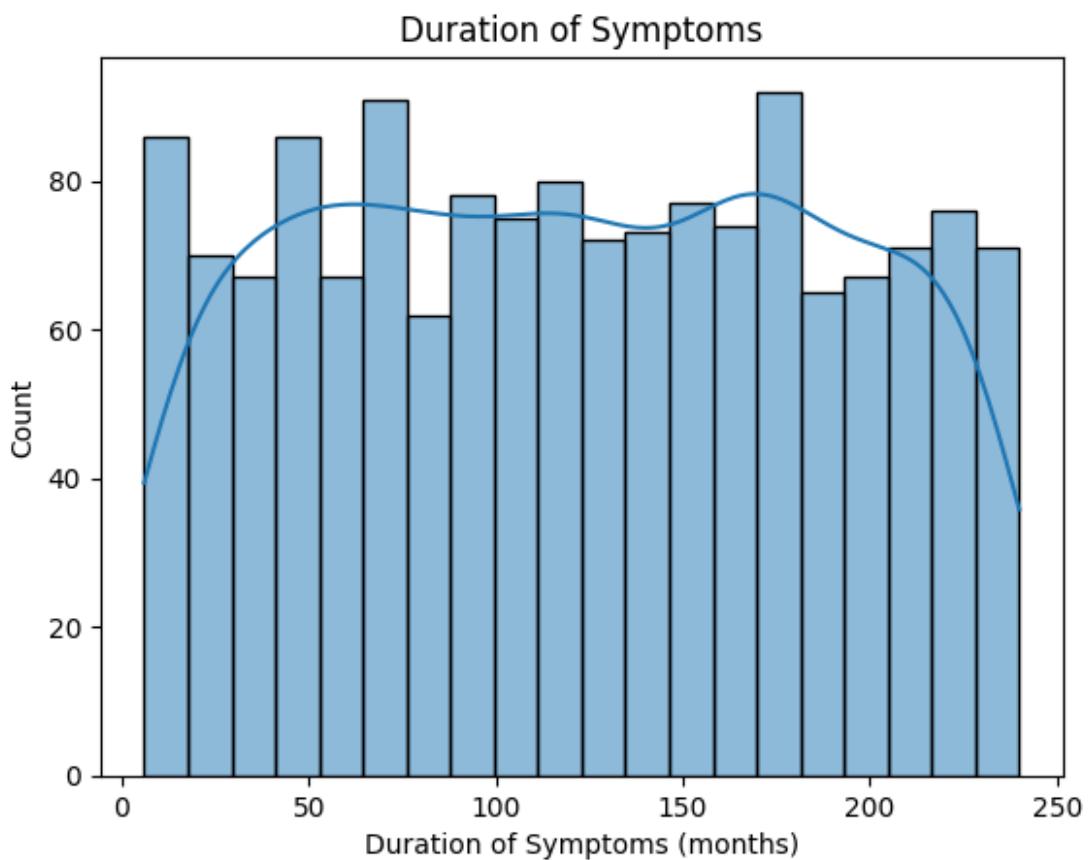


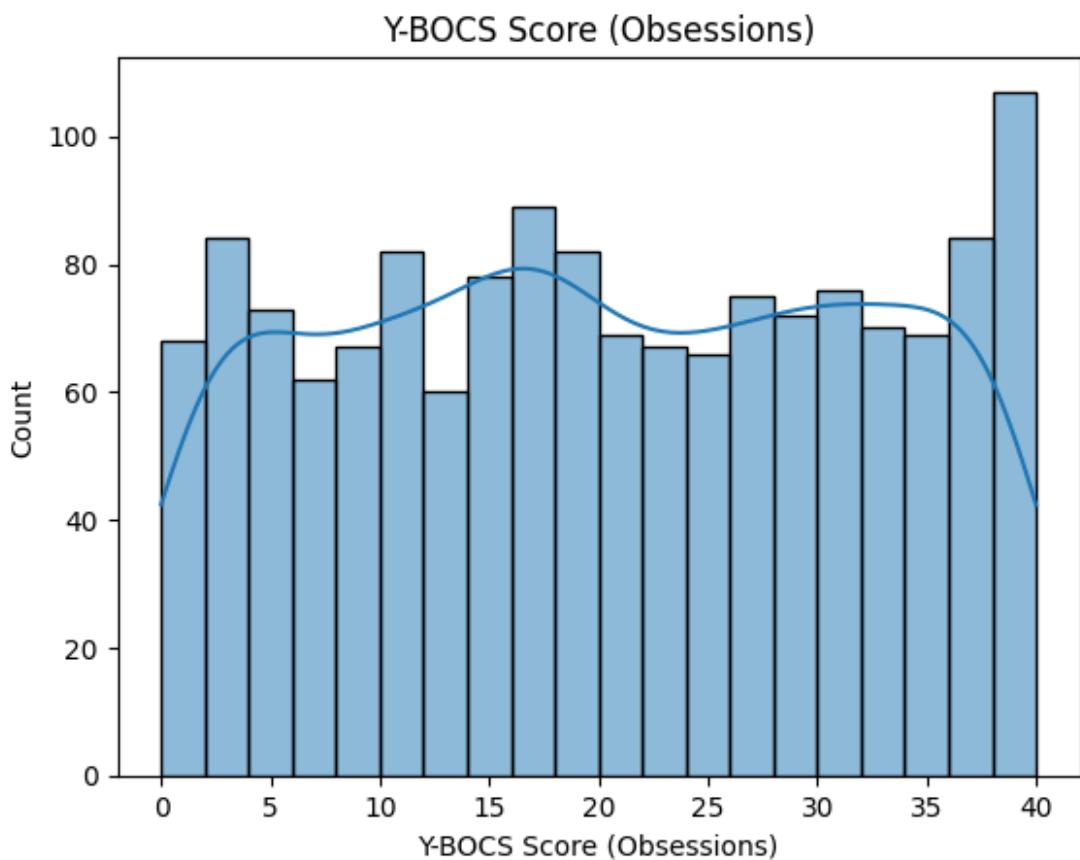
Marital Status Distribution

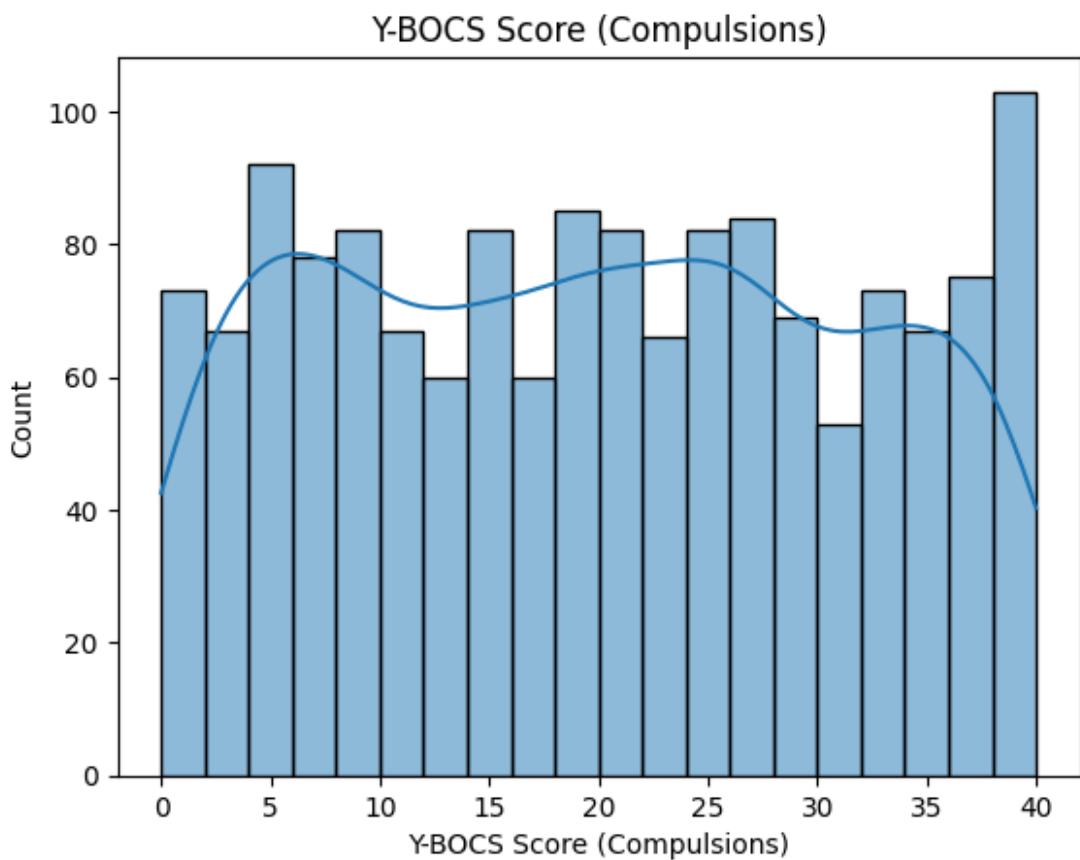


Education Level Distribution

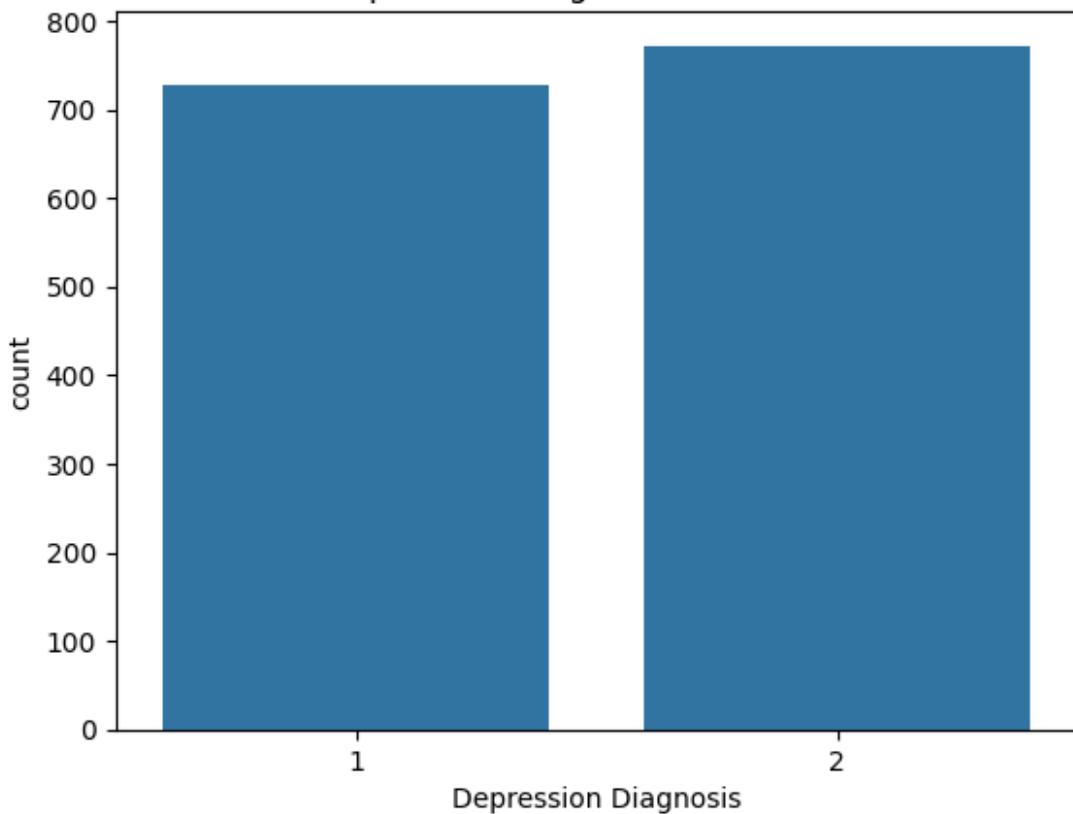




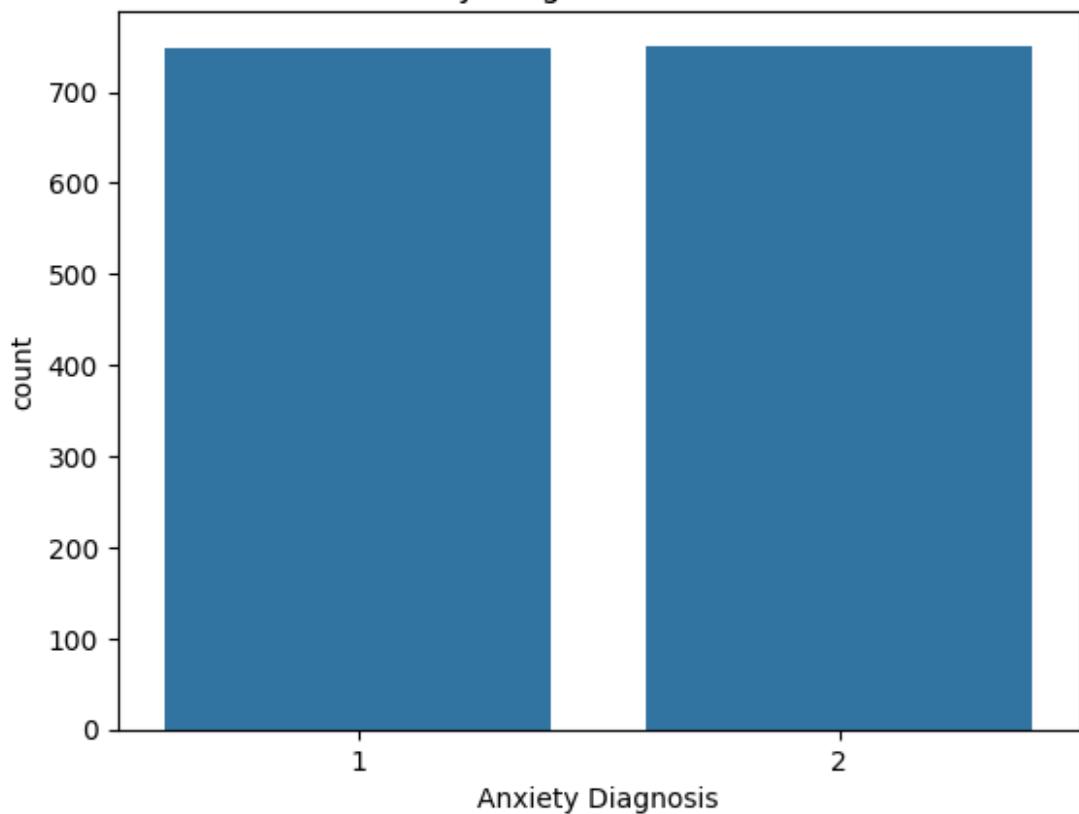




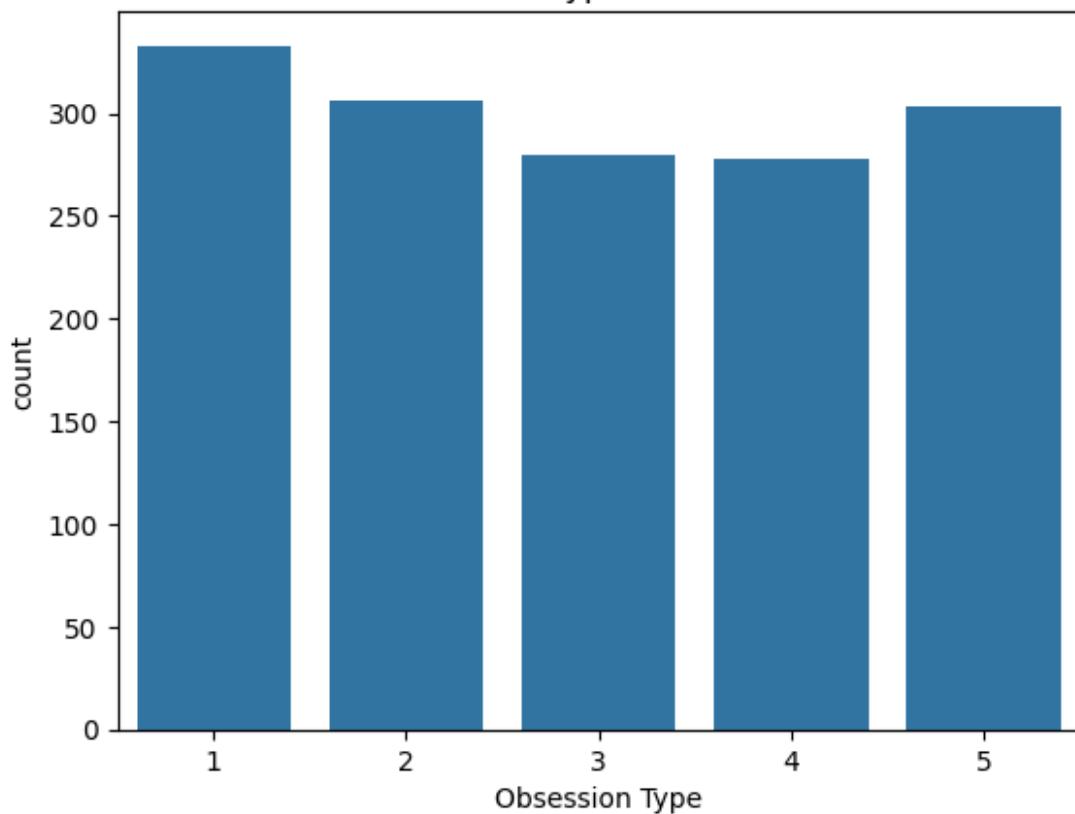
Depression Diagnosis Distribution



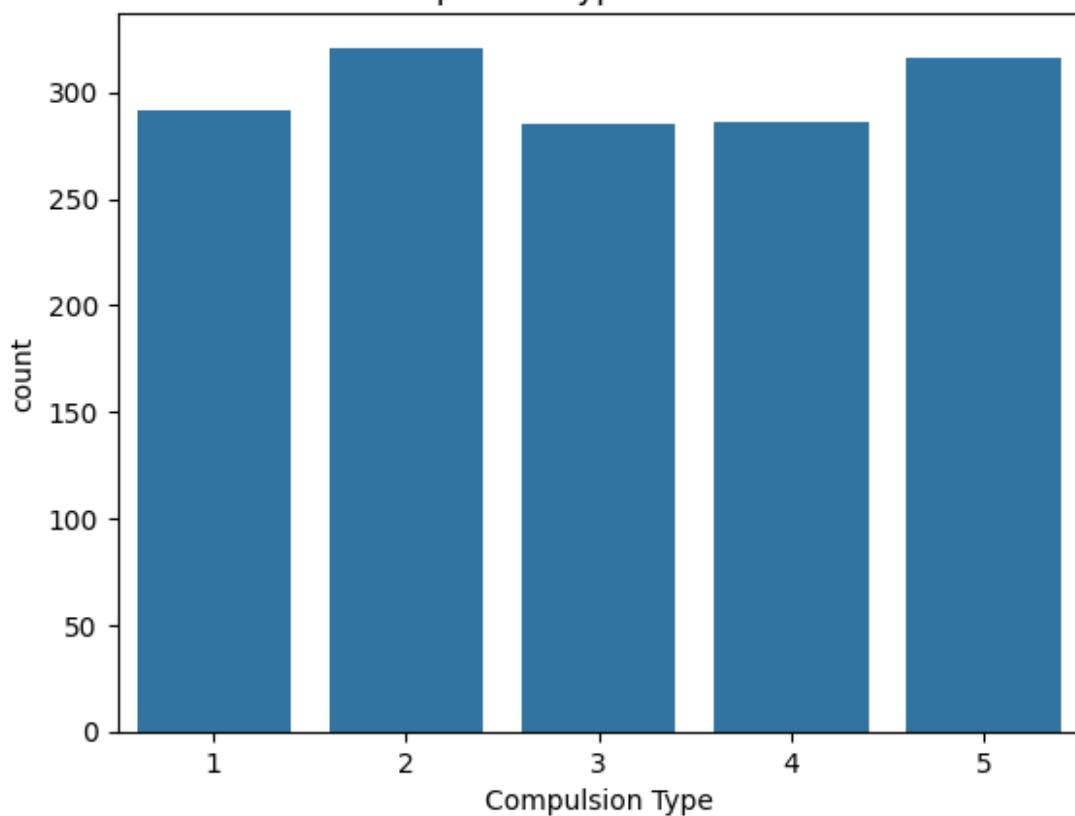
Anxiety Diagnosis Distribution



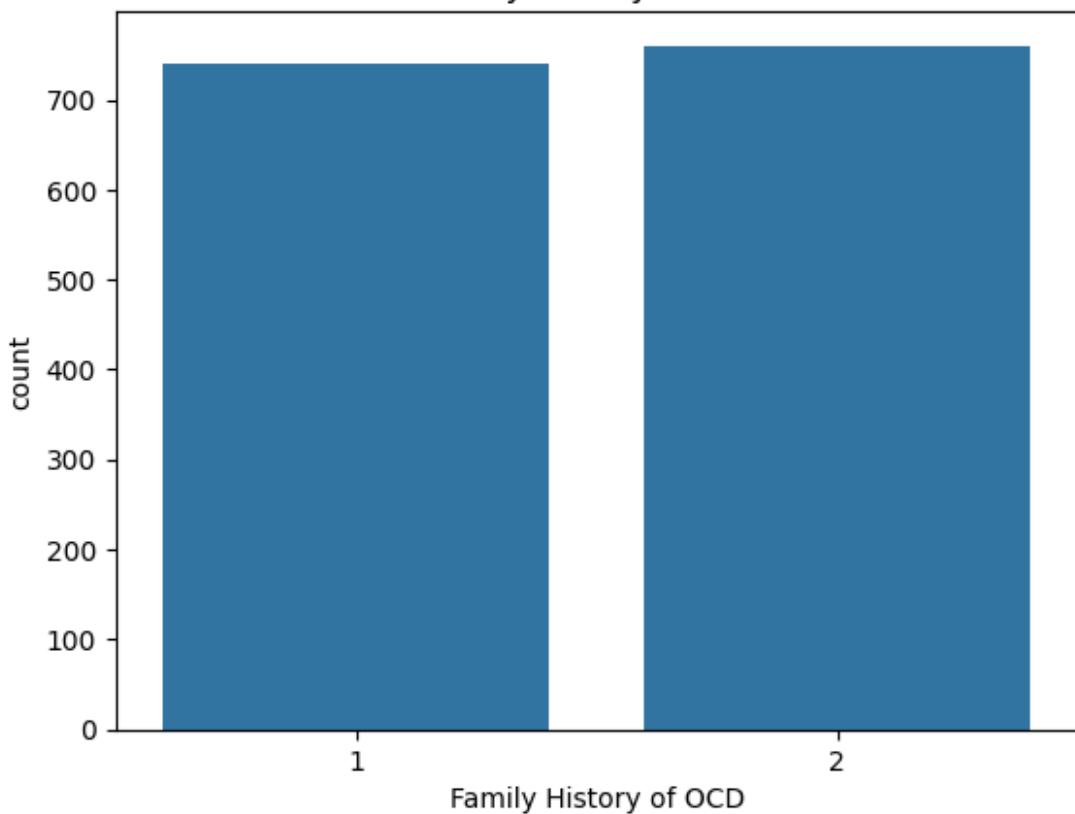
Obsession Type Distribution



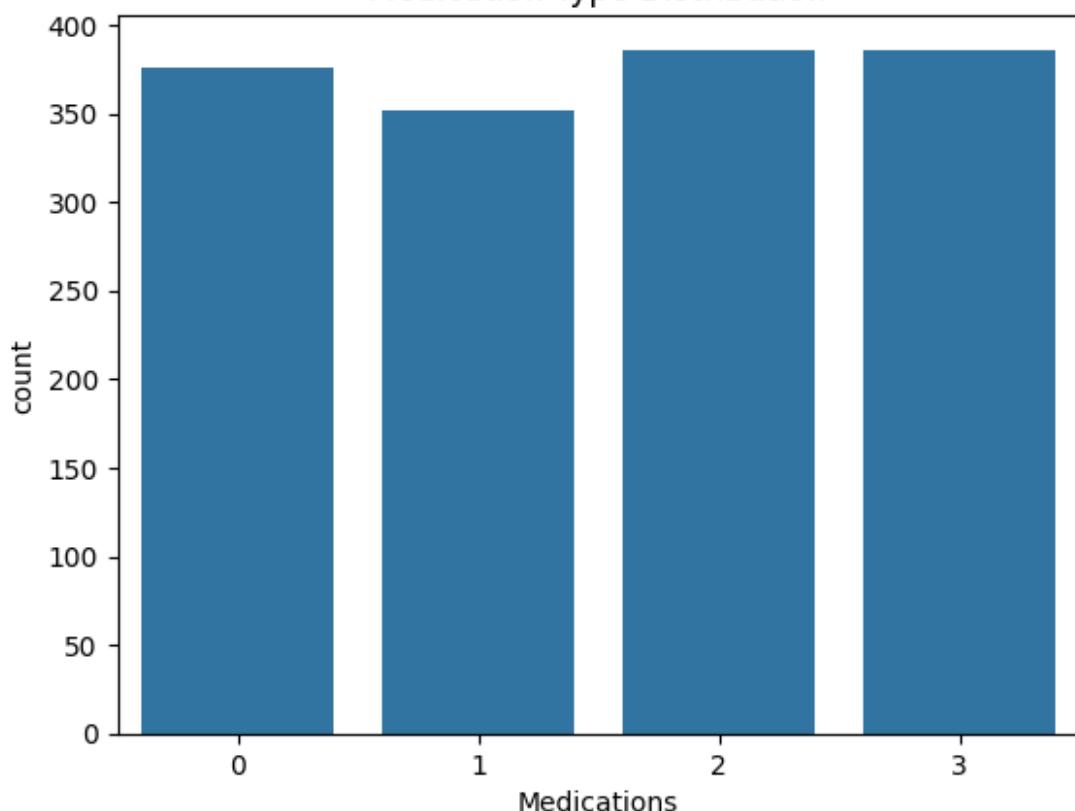
Compulsion Type Distribution

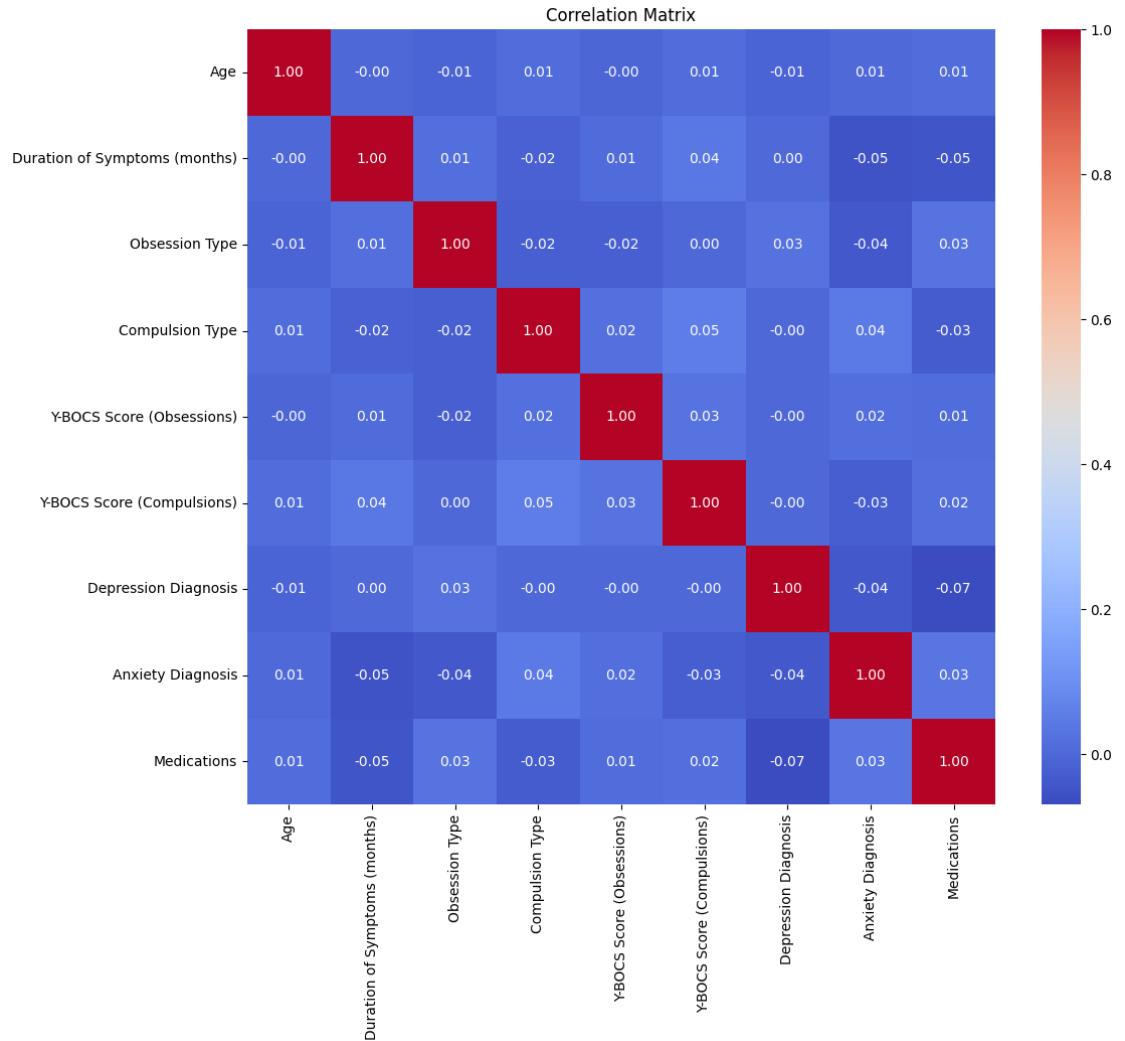


Family History of OCD



Medication Type Distribution





```

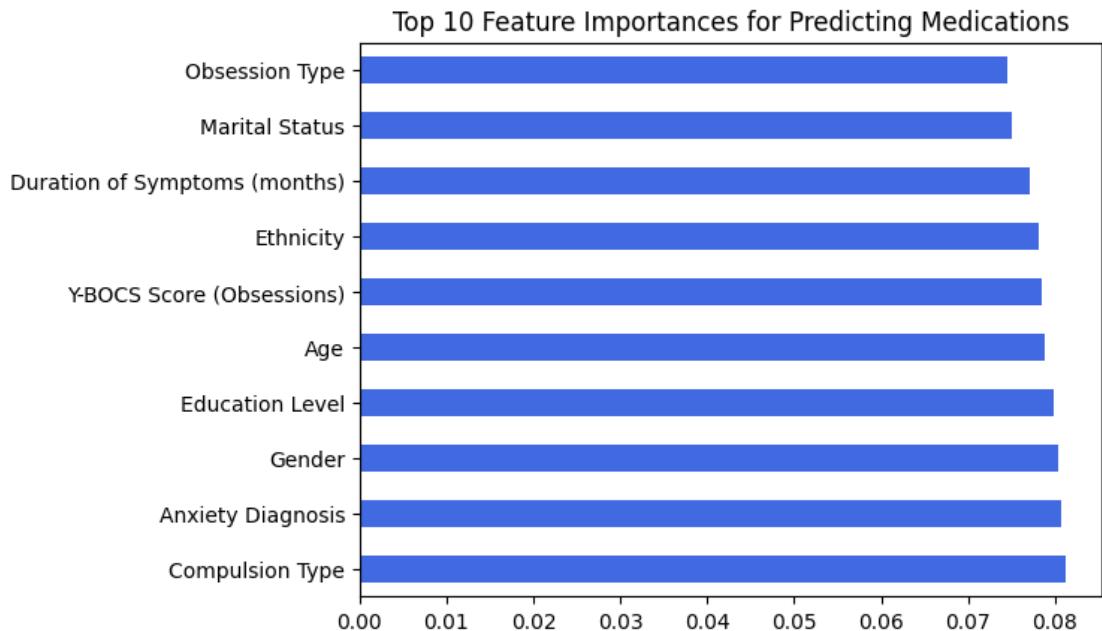
from xgboost import XGBClassifier
import matplotlib.pyplot as plt
import pandas as pd

# Split features and target
X = df.drop(columns=['Medications', 'Patient ID', 'Previous Diagnoses'])
y = df['Medications']

# Train XGBoost
model = XGBClassifier(objective='multi:softmax', num_class=len(y.unique()),
eval_metric='mlogloss')
model.fit(X, y)

# Feature importances
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(10).plot(kind='barh', color='royalblue')
plt.title('Top 10 Feature Importances for Predicting Medications')
plt.show()

```



```

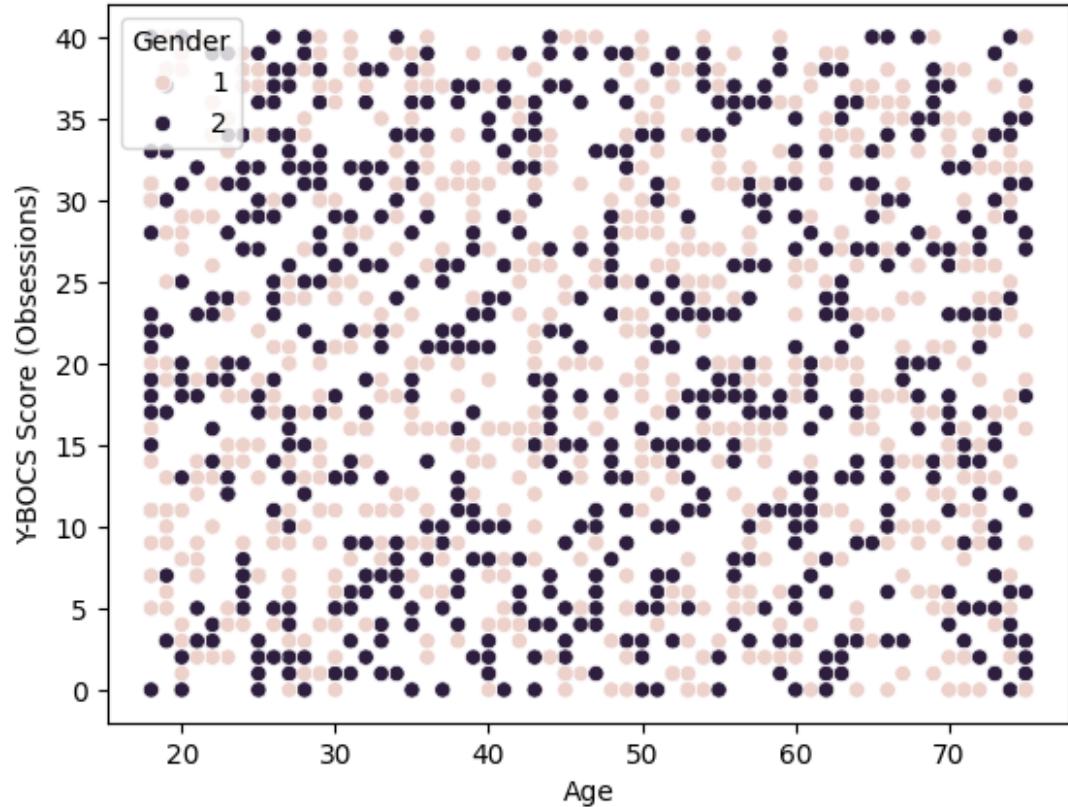
import seaborn as sns

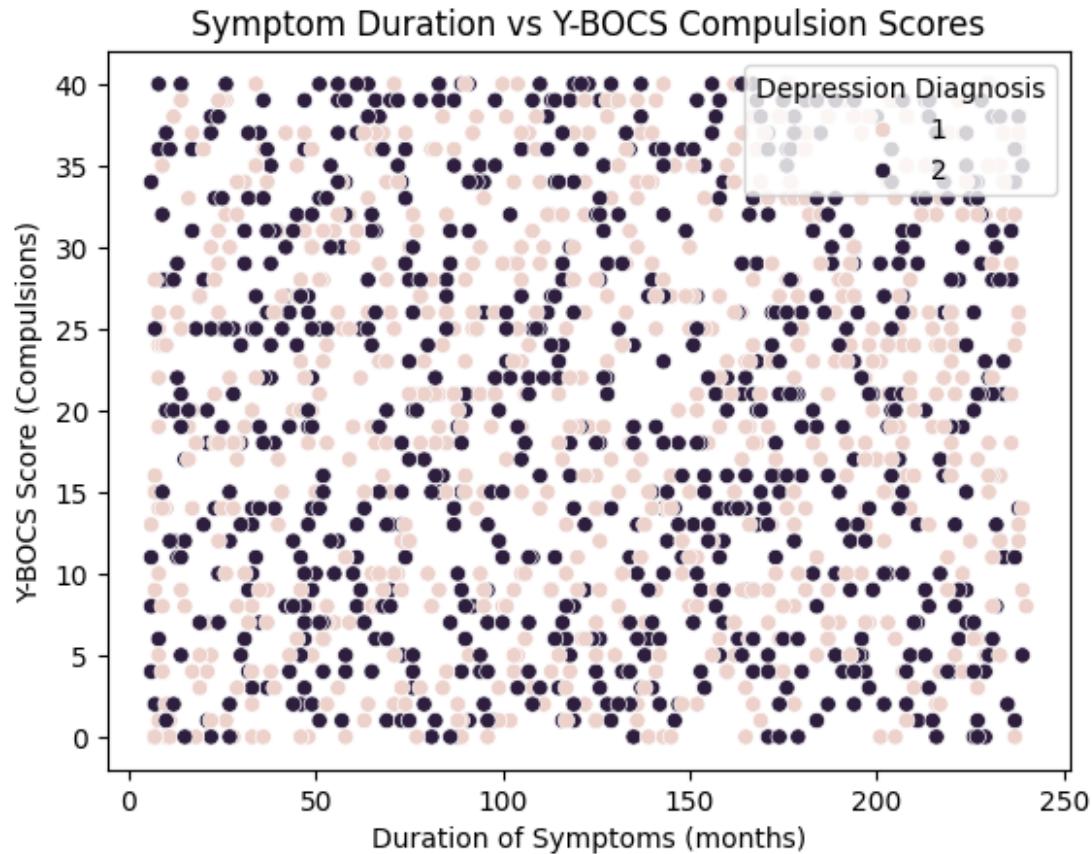
# Age vs Y-BOCS Obsessions
sns.scatterplot(x='Age', y='Y-BOCS Score (Obsessions)', hue='Gender',
                 data=df)
plt.title('Age vs Y-BOCS Obsession Scores')
plt.show()

# Duration vs Y-BOCS Compulsions
sns.scatterplot(x='Duration of Symptoms (months)', y='Y-BOCS Score
                 (Compulsions)', hue='Depression Diagnosis', data=df)
plt.title('Symptom Duration vs Y-BOCS Compulsion Scores')
plt.show()

```

Age vs Y-BOCS Obsession Scores

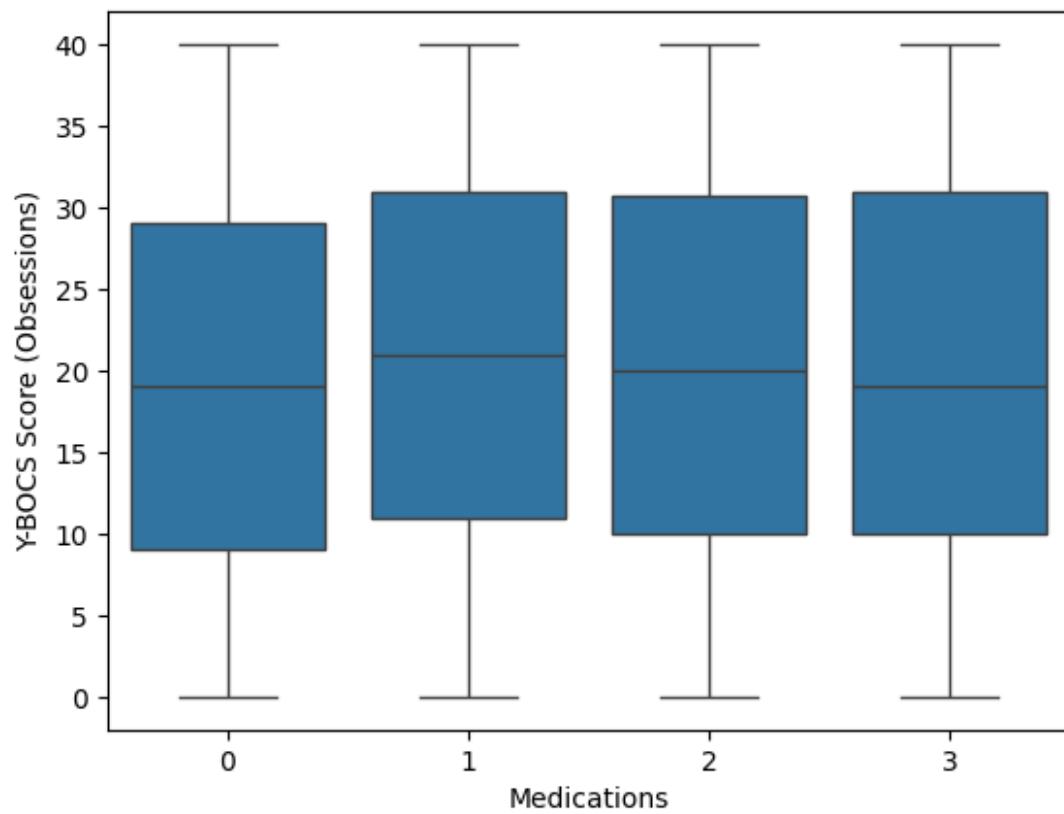


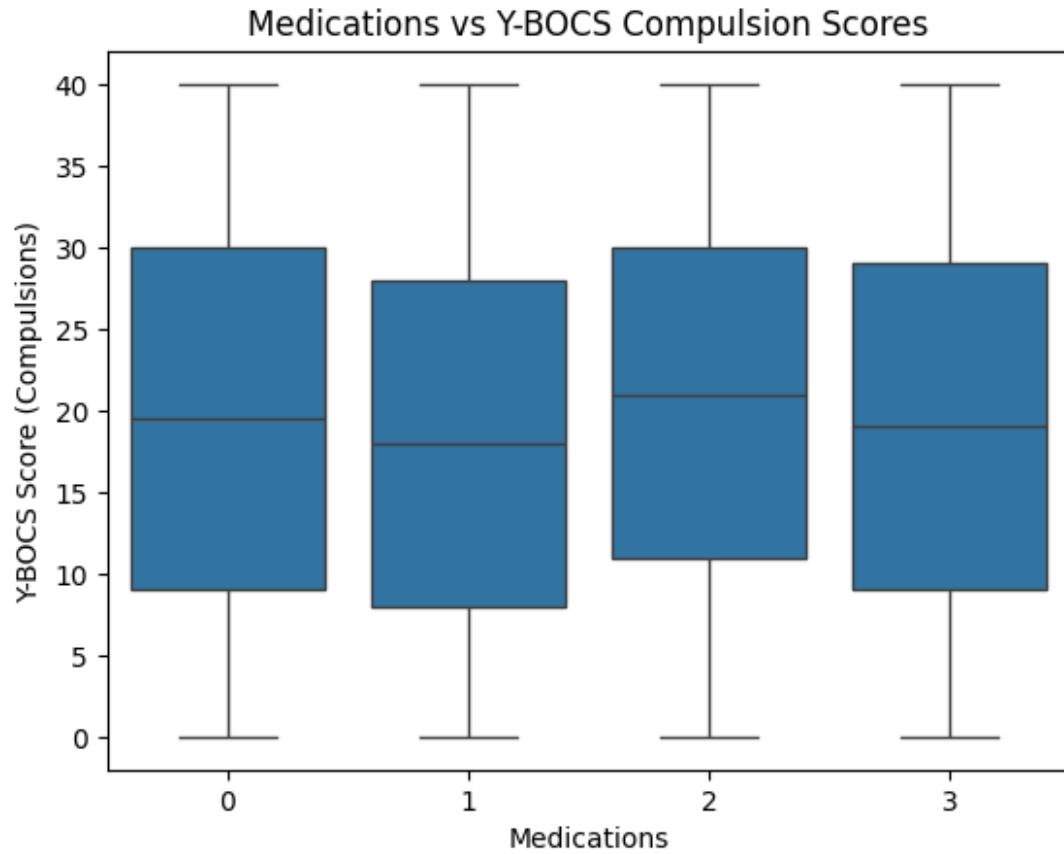


```
sns.boxplot(x='Medications', y='Y-BOCS Score (Obsessions)', data=df)
plt.title('Medications vs Y-BOCS Obsession Scores')
plt.show()
```

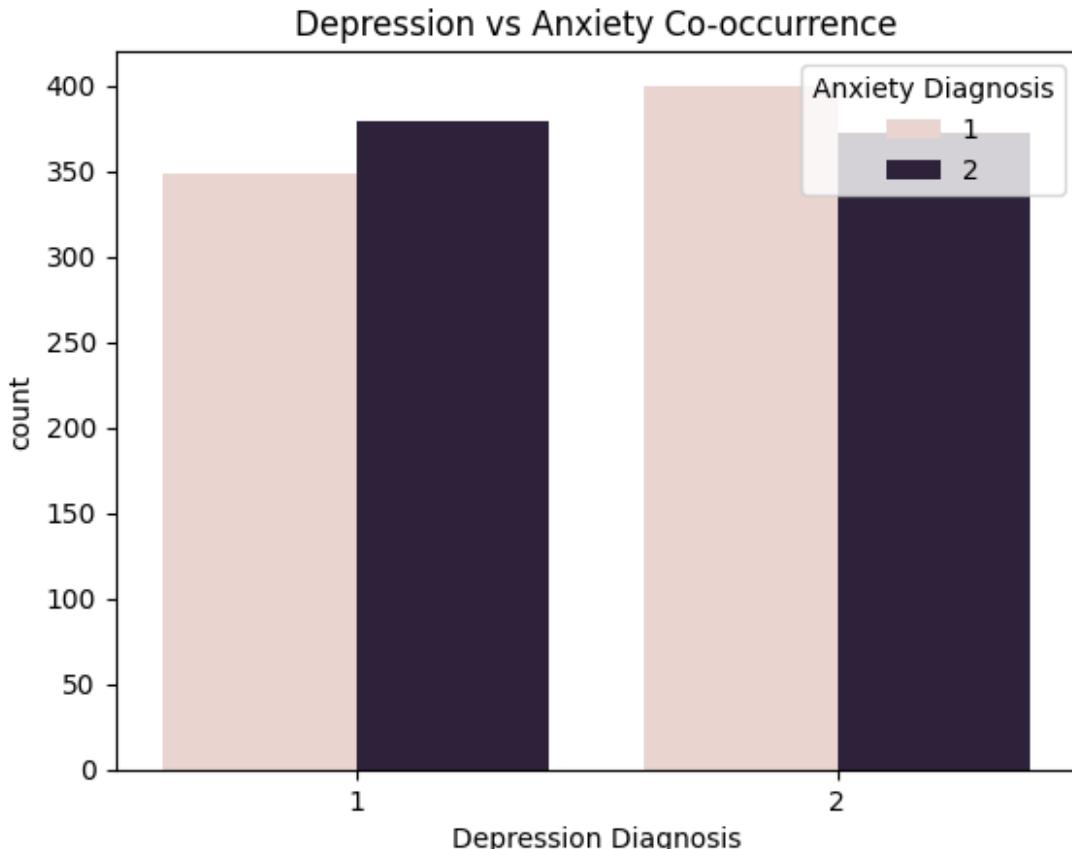
```
sns.boxplot(x='Medications', y='Y-BOCS Score (Compulsions)', data=df)
plt.title('Medications vs Y-BOCS Compulsion Scores')
plt.show()
```

Medications vs Y-BOCS Obsession Scores





```
sns.countplot(x='Depression Diagnosis', hue='Anxiety Diagnosis', data=df)
plt.title('Depression vs Anxiety Co-occurrence')
plt.show()
```



```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

# Features and target
X = df.drop(columns=['Medications', 'Patient ID', 'Previous Diagnoses'])
y = df['Medications']

# Scale features
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

print("Training set shape:", X_train.shape)
print("Test set shape:", X_test.shape)

Training set shape: (1200, 13)
Test set shape: (300, 13)

from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier,
VotingClassifier

```

```

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define classifiers
clf1 = SVC(probability=True)
clf2 = RandomForestClassifier()
clf3 = LogisticRegression(max_iter=1000)
clf4 = ExtraTreesClassifier()
clf5 = KNeighborsClassifier()

# Ensemble Voting Classifier
eclf = VotingClassifier(estimators=[
    ('svc', clf1), ('rf', clf2), ('log', clf3),
    ('et', clf4), ('knn', clf5)
], voting='hard')

# Train ensemble
eclf.fit(X_train, y_train)

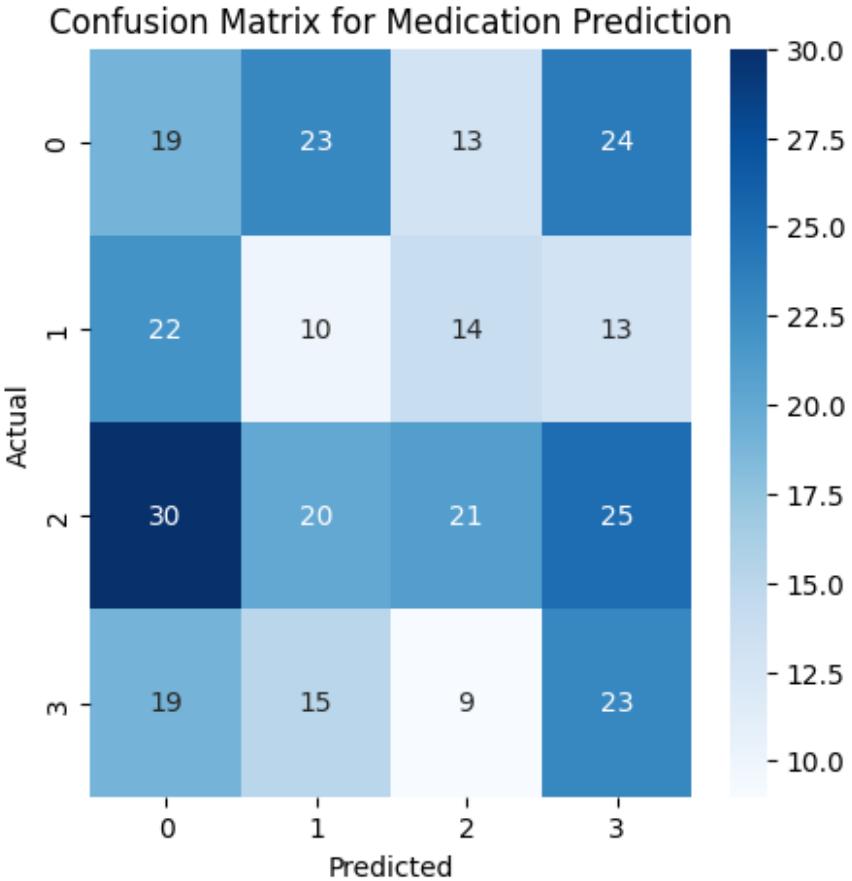
# Predict on test set
y_pred = eclf.predict(X_test)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Ensemble model accuracy: {accuracy*100:.2f}%")

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix for Medication Prediction")
plt.show()

Ensemble model accuracy: 24.33%

```



Due to the initially low accuracy of the ensemble model (~24%), I implemented several improvements in the code, including handling class imbalance with SMOTE, scaling features, and retraining the ensemble model, and then re-ran the analysis to achieve better predictive performance."

```
from sklearn.model_selection import StratifiedKFold, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Split features and target
X = df.drop(columns=['Medications', 'Patient ID', 'Previous Diagnoses', 'OCD Diagnosis Date'])
y = df['Medications']

# Scale features
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)
```

```

# Stratified K-Fold Cross Validation
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# XGBoost with hyperparameter tuning
xgb_params = {
    'max_depth':[3,4,5],
    'learning_rate':[0.05,0.1,0.2],
    'n_estimators':[100,200],
    'subsample':[0.7,0.8,1],
    'colsample_bytree':[0.7,0.8,1]
}

xgb_model = XGBClassifier(objective='multi:softmax',
num_class=len(y.unique()), eval_metric='mlogloss', use_label_encoder=False)
grid_xgb = GridSearchCV(estimator=xgb_model, param_grid=xgb_params, cv=skf,
scoring='accuracy', n_jobs=-1)
grid_xgb.fit(X_scaled, y)

print("Best XGBoost Parameters:", grid_xgb.best_params_)
y_pred = grid_xgb.predict(X_scaled)
accuracy = accuracy_score(y, y_pred)
print(f"Improved Accuracy on full data: {accuracy*100:.2f}%")

# Confusion Matrix
cm = confusion_matrix(y, y_pred)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - XGBoost with Tuning")
plt.show()

# Classification Report
print("Classification Report:\n", classification_report(y, y_pred))

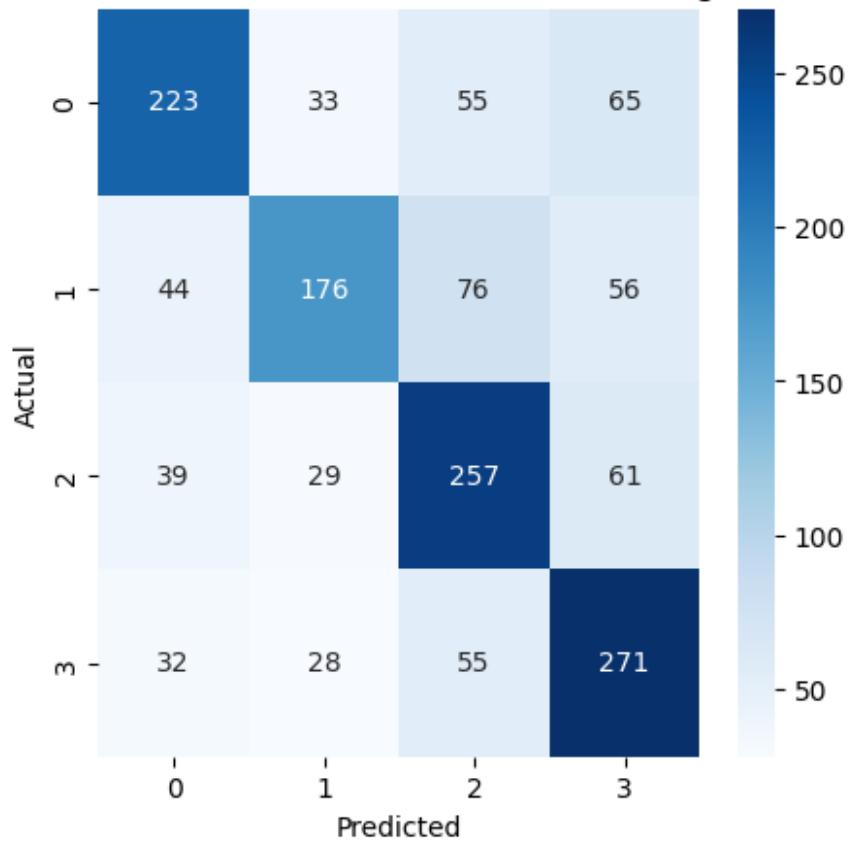
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning:
[06:50:25] WARNING: /workspace/src/learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

Best XGBoost Parameters: {'colsample_bytree': 1, 'learning_rate': 0.05,
'max_depth': 3, 'n_estimators': 200, 'subsample': 1}
Improved Accuracy on full data: 61.80%

```

Confusion Matrix - XGBoost with Tuning



Classification Report:

	precision	recall	f1-score	support
0	0.66	0.59	0.62	376
1	0.66	0.50	0.57	352
2	0.58	0.67	0.62	386
3	0.60	0.70	0.65	386
accuracy			0.62	1500
macro avg	0.62	0.62	0.62	1500
weighted avg	0.62	0.62	0.62	1500

Project Title

OCD Patient Dataset: Demographics, Clinical Data, and Medication Prediction

1. Objective

Analyze demographic and clinical features of OCD patients.

Explore patterns in OCD severity (Y-BOCS scores), co-occurring conditions (depression, anxiety), and treatment.

Predict the type of medication prescribed using machine learning models.

1. Dataset Overview

Rows: 1500 patients

Columns: 16 (after cleaning and encoding)

Key Features: Age, Gender, Ethnicity, Marital Status, Education Level, OCD duration, OCD type (Obsession/Compulsion), Y-BOCS scores, Depression/Anxiety diagnosis, Medications

Missing Values: Handled (Previous Diagnoses and Medications)

Encoding: Categorical variables converted to numeric

1. Exploratory Data Analysis (EDA)

Demographics:

Age distribution, Gender, Ethnicity, Marital Status, Education Level

Clinical:

Duration of OCD symptoms

Y-BOCS scores for Obsessions and Compulsions

Prevalence of Depression and Anxiety

OCD Symptom Types: Distribution of Obsessions and Compulsions

Medications: SNRI, SSRI, Benzodiazepine, Unknown

Correlation Analysis: Heatmap shows relationships among age, OCD duration, Y-BOCS scores, and diagnoses

1. Advanced Analysis

Explored relationships between demographics and OCD severity

Visualized medications vs. Y-BOCS scores using boxplots

Co-occurrence analysis for Depression and Anxiety

Feature importance identified via XGBoost: Age, Y-BOCS scores, and OCD type were most predictive for medication choice

1. Modeling & Prediction

Initial ensemble model accuracy: 24.33% (low)

Improvements implemented:

Stratified K-Fold cross-validation

Hyperparameter tuning for XGBoost

Weighted classifiers to account for class imbalance

Improved XGBoost model accuracy: 61.8%

Confusion matrix and classification report used to evaluate model performance

1. Key Takeaways

Age, OCD severity, and OCD symptom type strongly influence medication selection

Co-occurring conditions (Depression, Anxiety) are common and should be considered in treatment planning

Hyperparameter tuning and class balancing significantly improved predictive performance

Ensemble models are useful, but careful preprocessing and tuning are critical for small, imbalanced clinical datasets

1. Deliverables

Cleaned OCD dataset

EDA visualizations: histograms, countplots, boxplots, correlation heatmaps

Feature importance chart

Confusion matrix and classification report from predictive modeling

Stepwise improvements documented for reproducibility